

PRD: “ReadMe” — Local AI Reading Assistant

1. Purpose and Vision

ReadMe is a **self-hosted reading application** that converts digital books and documents (**.pdf**, **.epub**, **.txt**, **.docx**) into natural-sounding speech and intelligent summaries. It is designed for **personal, local use** — running on your PC — with **optional cloud compute on Heroku** for resource-intensive AI tasks.

The intent is to build a **Speechify-like desktop utility** that prioritizes:

- privacy and ownership of data,
 - local responsiveness,
 - zero external dependencies except for API calls you explicitly configure.
-

2. Functional Overview

| Feature | Description |
|---------------------------|--|
| File Input | Open .pdf , .epub , .txt , .docx files from local drive. |
| Text Extraction | Parse text and structure (pages, chapters, paragraphs). |
| Speech Synthesis | Convert text to speech using AI voices (via local Coqui-TTS or OpenAI API through Heroku). |
| Playback Controls | Play, pause, seek, adjust playback speed, repeat paragraph, or skip. |
| Library Management | Maintain local list of previously opened books with reading progress. |
| Bookmarks / Notes | Save reading position, annotations, and highlights locally. |
| AI Summarization | Summarize selected text or chapters (processed on Heroku). |

| | |
|--------------------------|--|
| Offline Mode | Fully functional offline reading (except for remote AI services). |
| Minimal Interface | Lightweight desktop GUI with file explorer, playback, and settings panels. |

3. System Architecture

Overview

- **Frontend + Local App:** Electron + React (desktop client).
- **Local Backend:** Python FastAPI service running on localhost for parsing, TTS control, and storage.
- **Remote Backend (Heroku):** Handles large AI workloads like text summarization, OpenAI TTS, or Whisper transcription.

Architecture Layers

A. Local Components

| Layer | Tech | Purpose |
|-----------------------------|---------------------------------|--|
| Desktop Shell | Electron + React | GUI for book browsing, playback, and settings. |
| Local API Server | FastAPI (Python) | Hosts endpoints for local logic and file parsing. |
| Database | SQLite | Store library metadata, progress, notes, and settings. |
| File Parsing | PyMuPDF / pdfplumber / ebooklib | Extracts text + metadata from documents. |
| Local TTS (optional) | Coqui-TTS | Perform on-device text-to-speech when available. |
| Audio Engine | PyAudio / Web Audio API | Stream audio to local output. |

B. Heroku Cloud Components

| Service | Tech | Role |
|---------|------|------|
|---------|------|------|

| | | |
|-----------------------------------|--|--|
| Heroku Web App | FastAPI (Python) / Flask | Receives text from local app, returns summarized or TTS-generated audio. |
| AI Layer | OpenAI API, Whisper, Transformers | Summarization, translation, advanced speech synthesis. |
| Heroku Postgres (optional) | Used only for caching API results or analytics if desired. | |

4. Data Flow

1. User Opens File

- Electron app reads file metadata and sends to local FastAPI.
- FastAPI extracts text and structures it (JSON: `{chapter, page, paragraph}`).

2. Speech Generation

- For small jobs → Coqui-TTS locally.
- For large jobs → sends text to Heroku FastAPI endpoint (`/generate_audio`).
- Heroku service calls OpenAI TTS or another model, streams back `.wav/.mp3`.

3. Playback

- Local player streams from local disk or directly from Heroku audio stream.

4. Summarization

- Selected text sent to Heroku `/summarize` endpoint.
- Response returned to desktop app and cached locally.

5. Persistence

- SQLite DB stores:

- file paths
 - read progress
 - bookmarks & annotations
 - TTS cache locations
-

5. Software Stack

Local (Desktop)

| Category | Tech Stack |
|------------------------|---|
| GUI | Electron + React + TailwindCSS |
| Local API | FastAPI (Python 3.12) |
| Parsing | <code>pdfplumber</code> , <code>ebooklib</code> , <code>PyMuPDF</code> , <code>docx2txt</code> |
| TTS (optional local) | Coqui-TTS |
| Storage | SQLite |
| Audio Playback | PyAudio or HTML5 Audio via Electron |
| Environment Management | Poetry / Pipenv |
| Packaging | Electron Builder |

Heroku (Cloud)

| Layer | Tech |
|--------------|---|
| Web Server | FastAPI / Flask |
| Task Queue | Celery + Redis (for longer TTS jobs) |
| AI Libraries | OpenAI API, HuggingFace Transformers, Whisper |
| Storage | Local ephemeral or S3-compatible bucket |

Deployment Heroku Pipelines (Staging / Prod)

Security API key authentication (personal key only)

6. Local–Cloud Integration

Local Endpoint:

```
POST /api/tts
{
  "text": "...",
  "voice": "rich-voice-1",
  "mode": "cloud" | "local"
}
```

Cloud Endpoint (Heroku):

```
POST /api/v1/tts
{
  "text": "...",
  "model": "gpt-4o-tts",
  "voice": "alloy"
}
```

Heroku Returns:

- Streamed or pre-generated .mp3
 - Response metadata (duration, sample_rate, voice_used)
-

7. Security Model

- Local FastAPI only binds to localhost:5000 (not exposed publicly).

- All cloud interactions use HTTPS with personal bearer token.
 - No third-party user accounts or telemetry.
 - Heroku stores no user data — ephemeral compute only.
-

8. Performance and Scalability

| Operation | Location | Target Performance |
|------------------------|--------------|---------------------|
| Parsing a 200-page PDF | Local | < 5 seconds |
| Local TTS (Coqui) | Local | 1x real-time |
| Cloud TTS (Heroku) | Remote | < 2 sec latency |
| Summarization (GPT-4o) | Remote | < 8 sec per chapter |
| DB read/write | Local SQLite | instantaneous |

9. Deployment & DevOps

| Area | Tool |
|------------------------------|---|
| Local App Packaging | Electron Builder (creates <code>.app</code> / <code>.exe</code>) |
| Backend Hosting | Heroku Standard Dyno |
| Database Migration | Alembic (for Heroku Postgres if used) |
| Version Control | Git + GitHub |
| CI/CD | GitHub Actions → Auto-deploy to Heroku |
| Local Runtime | Docker Compose (optional) |
| Logs & Monitoring | Heroku Logs + Sentry (local + cloud) |

10. Roadmap

| Phase | Deliverables |
|----------------------|--|
| Phase 1 (MVP) | Electron app + local FastAPI + PDF/EPUB parsing + basic playback |
| Phase 2 | Integrate Heroku TTS + caching + voice selector |
| Phase 3 | Add summarization endpoint + local note system |
| Phase 4 | Coqui-TTS local fallback + offline speech synthesis |
| Phase 5 | UI polish + exportable audio + hotkey navigation |

11. Directory Layout (Proposed)

```
readme-app/
├── frontend/                      # Electron + React
│   ├── src/
│   │   ├── components/
│   │   ├── pages/
│   │   └── store/
│   └── package.json
├── backend/                         # FastAPI local API
│   ├── main.py
│   ├── tts/
│   │   ├── coqui.py
│   │   └── openai_cloud.py
│   └── parsers/
│       ├── pdf_parser.py
│       ├── epub_parser.py
│       └── doc_parser.py
└── cloud/                           # Heroku FastAPI entry
    ├── app.py
    ├── summarize.py
    ├── tts_openai.py
    └── requirements.txt
└── db/
    └── readme.db
└── config/
    ├── settings.yaml
    └── secrets.env
```

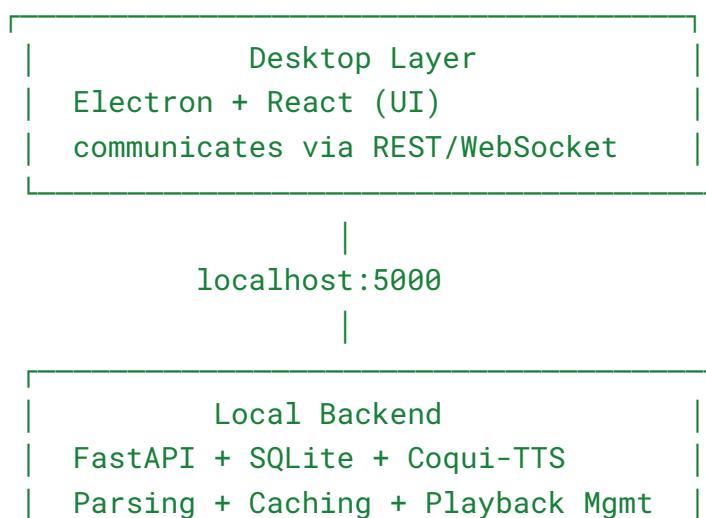
12. Technologies Summary

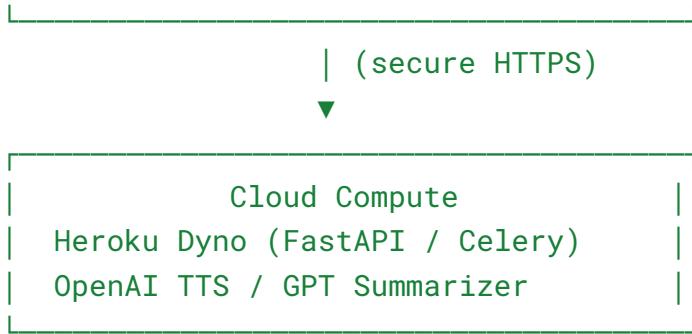
| Layer | Stack |
|------------------|---|
| Frontend | Electron + React + TailwindCSS |
| Backend (Local) | FastAPI, PyMuPDF, pdfplumber, Coqui-TTS, SQLite |
| Backend (Heroku) | FastAPI, OpenAI API, Celery, Redis |
| Infrastructure | Docker, GitHub Actions, Heroku Dynos |
| Security | API key, localhost-only exposure, TLS |
| Audio Pipeline | Local caching, async streaming |

Technical Specification (TS): “ReadMe” – Local AI Reading Assistant

1. System Topology

Architecture Overview





2. Core Data Structures

2.1 Book Metadata

```
class Book(BaseModel):
    id: str                      # UUID
    title: str
    author: Optional[str]
    filepath: str                 # Absolute local path
    filetype: str                 # pdf | epub | txt | docx
    total_pages: int
    last_read_page: int = 0
    last_accessed: datetime
    cover_image_path: Optional[str]
    text_cache_path: Optional[str]
```

2.2 Annotation / Note

```
class Annotation(BaseModel):
    id: str
    book_id: str
    page_number: int
    selection_text: str
    note_text: str
    timestamp: datetime
```

2.3 Audio Job (Local/Remote)

```
class AudioJob(BaseModel):
    id: str
    book_id: str
    chapter: str
    status: str          # queued | processing | done | failed
    engine: str          # local | cloud
    voice: str
    text_path: str
    audio_path: Optional[str]
    created_at: datetime
```

3. Database Schema (SQLite)

```
CREATE TABLE books (
    id TEXT PRIMARY KEY,
    title TEXT,
    author TEXT,
    filepath TEXT,
    filetype TEXT,
    total_pages INTEGER,
    last_read_page INTEGER,
    last_accessed DATETIME,
    cover_image_path TEXT,
    text_cache_path TEXT
);

CREATE TABLE annotations (
    id TEXT PRIMARY KEY,
    book_id TEXT,
    page_number INTEGER,
    selection_text TEXT,
    note_text TEXT,
    timestamp DATETIME,
    FOREIGN KEY(book_id) REFERENCES books(id)
);

CREATE TABLE audio_jobs (
```

```

    id TEXT PRIMARY KEY,
    book_id TEXT,
    chapter TEXT,
    status TEXT,
    engine TEXT,
    voice TEXT,
    text_path TEXT,
    audio_path TEXT,
    created_at DATETIME,
    FOREIGN KEY(book_id) REFERENCES books(id)
);

```

4. API Specification

4.1 Local FastAPI (localhost:5000)

| Method | Endpoint | Description |
|--------|----------------------------|---|
| POST | /api/books/import | Upload & parse .pdf, .epub, etc. Returns metadata. |
| GET | /api/books | List all local books with metadata. |
| GET | /api/books/{id} | Get details, progress, and cached content. |
| POST | /api/books/{id}/parse | Force re-parse a file. |
| POST | /api/tts | Generate audio (mode: <code>local</code> or <code>cloud</code>). |
| GET | /api/audio/{id}/stream | Stream cached .wav/.mp3 to player. |
| POST | /api/annotate | Add a note/bookmark. |
| GET | /api/annotations/{book_id} | Retrieve annotations. |
| POST | /api/summarize | Send selected text to Heroku summarization service. |
| GET | /api/settings | Retrieve system config (voices, cache paths). |

PUT /api/settings Update local preferences.

4.2 Cloud Heroku FastAPI (<https://readme-ai.herokuapp.com>)

| Method | Endpoint | Description |
|--------|-------------------|--|
| POST | /api/v1/tts | Convert text → audio using OpenAI TTS or another engine. |
| POST | /api/v1/summarize | Summarize a paragraph or chapter. |
| POST | /api/v1/translate | (Optional) Translate selected text. |
| GET | /api/v1/health | Health check. |

4.3 Example Payloads

Local → Cloud TTS Request

```
POST /api/v1/tts
{
  "text": "Chapter one begins...",
  "model": "gpt-4o-tts",
  "voice": "alloy"
}
```

Response

```
{
  "job_id": "a2b1f...",
  "duration": 32.6,
  "sample_rate": 44100,
  "audio_url": "https://readme-ai.herokuapp.com/files/a2b1f.mp3"
}
```

5. Processing Pipelines

5.1 File Parsing (Local)

1. Electron selects file → POST `/api/books/import`
2. FastAPI determines parser based on file extension

Extracted text stored in JSON cache:

`cache/books/<book_id>/text.json`

- 3.
4. DB entry created.

5.2 Text-to-Speech

- **Local Mode:** Coqui-TTS engine generates audio to cache directory.
- **Cloud Mode:** Text sent to Heroku → processed via OpenAI API → audio URL returned → downloaded + cached locally.

5.3 Summarization

1. Electron highlights passage → POST `/api/summarize`
2. FastAPI forwards to Heroku `/api/v1/summarize`
3. GPT-4o generates abstract → result cached in SQLite.

6. Heroku Cloud Services

6.1 Heroku App Structure

```
cloud/  
└── app.py  
└── routers/
```

```
|   └── tts.py
|   └── summarize.py
|── services/
|   ├── openai_client.py
|   └── audio_utils.py
└── Procfile
```

6.2 Example Heroku Procfile

```
web: gunicorn app:app --workers=1 --timeout 120
worker: celery -A tasks.celery_app worker --loglevel=info
```

6.3 Celery Task Example

```
@app.task
def generate_audio_task(text, voice="alloy"):
    audio = openai.audio.speech.create(
        model="gpt-4o-mini-tts",
        voice=voice,
        input=text
    )
    filename = f"{uuid.uuid4()}.mp3"
    with open(f"/tmp/{filename}", "wb") as f:
        f.write(audio.read())
    return f"/tmp/{filename}"
```

7. Desktop Application (Electron + React)

7.1 UI Layout

| Component | Purpose |
|-------------|--|
| Sidebar | Library (Book list, progress bar) |
| Reader Pane | Paginated text view, highlighting, annotations |
| Player Bar | Play/pause, speed, seek, voice selector |

| | |
|--------------------------------|---|
| Settings Panel | Choose local vs cloud mode, manage voices |
| Console Panel (Dev) | Display FastAPI logs, status updates |

7.2 IPC Communication

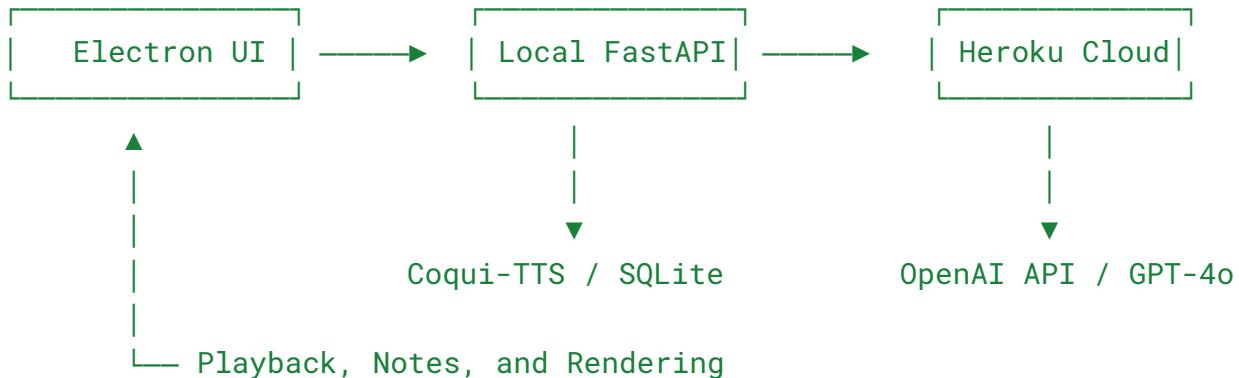
Electron bridges frontend ↔ Python backend:

```
ipcRenderer.invoke('fetch-books')
ipcRenderer.invoke('start-tts', { text, mode: 'local' })
ipcRenderer.invoke('summarize-text', { text })
```

Electron executes Python server using:

```
const pyProc = spawn('python', ['backend/main.py'])
```

8. Data Flow Diagram



9. Configuration and Secrets

```
config/settings.yaml
```

```
tts_default: "local"
voices:
  - "coqui_en"
```

```
- "openai_alloy"
heroku_api_url: "https://readme-ai.herokuapp.com"
api_key: !ENV ${HEROKU_API_KEY}
cache_dir: "./cache"

config/secrets.env

OPENAI_API_KEY=sk-...
HEROKU_API_KEY=...
```

10. Performance Targets

| Function | Target |
|------------------------|-----------------|
| PDF parse (200 pages) | < 5 sec |
| Local TTS latency | < 1.5x realtime |
| Cloud TTS round-trip | < 2 sec |
| Summarization response | < 8 sec |
| App cold start | < 3 sec |
| Memory footprint | < 600 MB |

11. Error Handling and Recovery

| Scenario | Mitigation |
|----------------|--------------------------------|
| Heroku offline | Fallback to local TTS |
| Parsing error | Retry with alternate parser |
| API timeout | Queue and retry async |
| Disk full | Alert and clear cache |
| DB locked | Rollback + exponential backoff |

12. Testing Strategy

| Level | Tools | Purpose |
|-------------|-------------------------|-----------------------------|
| Unit Tests | Pytest, Jest | Core parsing, API endpoints |
| Integration | Postman, pytest-asyncio | Local ↔ Heroku endpoints |
| UI Tests | Cypress | GUI behavior |
| Load Tests | Locust | Parsing + TTS concurrency |
| E2E | Electron Test Kit | Full pipeline validation |

13. Deployment Guide Summary

Local Setup

```
git clone readme-app
cd backend && poetry install
poetry run unicorn main:app --reload
cd ../frontend && npm install && npm run electron-dev
```

1.

Heroku Setup

```
heroku create readme-ai
git push heroku main
heroku config:set OPENAI_API_KEY=sk-...
heroku ps:scale web=1 worker=1
```

2.

14. Security Summary

- **localhost-only** backend access.

- **Cloud auth** via static bearer key (personal).
 - **Encrypted storage** for annotations and cache (SQLCipher optional).
 - **All external calls HTTPS-only.**
 - **No analytics / telemetry / external tracking.**
-

15. Optional Enhancements

| Feature | Description |
|------------------------|---|
| Voice Cloning | Train voice using Coqui-VC and store locally. |
| Web Scraper | Import articles directly from URLs. |
| AI Q&A Mode | Ask contextual questions about chapters using GPT-4o. |
| Batch Export | Generate audiobook MP3 for entire book. |
| CLI Interface | Lightweight terminal mode for TTS playback. |

Cloud repo layout

```
readme-cloud/
├── app.py
├── settings.py
├── auth.py
└── routers/
    ├── tts.py
    └── summarize.py
└── services/
    ├── openai_tts.py
    └── openai_summarize.py
└── tests/
    ├── test_tts.py
    └── test_summarize.py
```

```
└─ requirements.txt  
└─ Procfile  
└─ runtime.txt
```

Environment variables (Heroku)

```
OPENAI_API_KEY=...  
API_BEARER=some-long-random-token          # your personal bearer  
for cloud auth  
AUDIO_FORMAT=mp3                            # mp3 | wav | ogg  
(optional)  
OPENAI_TTS_MODEL=gpt-4o-mini-tts            # default TTS model  
OPENAI_SUMMARY_MODEL=gpt-4.1-mini           # or another lightweight  
text model
```

Notes:

- The **Audio API** `audio/speech` endpoint with models like `gpt-4o-mini-tts` is the current OpenAI path for TTS. [OpenAI Platform+2](#)[OpenAI Platform+2](#)
 - Keep an eye on OpenAI's changelog/deprecations and the newer **Responses API** guidance for text to avoid breaking changes. [OpenAI Platform+2](#)[OpenAI Platform+2](#)
-

requirements.txt

```
fastapi==0.115.0  
uvicorn==0.30.6  
gunicorn==23.0.0  
pydantic==2.9.2  
httpx==0.27.2  
python-multipart==0.0.9
```

(The official `openai` SDK evolves; this skeleton uses `httpx` against stable REST endpoints.)

Procfile

```
web: gunicorn app:app --workers=1 --timeout=120
```

runtime.txt

```
python-3.12.6
```

settings.py

```
from pydantic import BaseSettings, Field

class Settings(BaseSettings):
    openai_api_key: str = Field(alias="OPENAI_API_KEY")
    api_bearer: str = Field(alias="API_BEARER")
    audio_format: str = Field(default="mp3", alias="AUDIO_FORMAT")
    openai_tts_model: str = Field(default="gpt-4o-mini-tts",
        alias="OPENAI_TTS_MODEL")
    openai_summary_model: str = Field(default="gpt-4.1-mini",
        alias="OPENAI_SUMMARY_MODEL")

    class Config:
        env_file = ".env"
        case_sensitive = True

settings = Settings()
```

auth.py (simple bearer auth)

```
from fastapi import Header, HTTPException, status
from settings import settings

def verify_bearer(authorization: str = Header(...,
convert_underscores=False)):
    # Expect: Authorization: Bearer <token>
    if not authorization or not
authorization.lower().startswith("bearer "):
        raise HTTPException(status_code=status.HTTP_401_UNAUTHORIZED,
detail="Missing bearer token")
    token = authorization.split(" ", 1)[1].strip()
    if token != settings.api_bearer:
        raise HTTPException(status_code=status.HTTP_403_FORBIDDEN,
detail="Invalid token")
    return True
```

services/openai_tts.py

```
import httpx
from typing import AsyncIterator
from settings import settings

# OpenAI Text-to-Speech via Audio API (audio/speech)
# Docs: Audio & TTS overview + models (e.g., gpt-4o-mini-tts).
# We'll stream bytes and forward as a generator.
# References: Audio & speech docs; Text-to-speech guide; Model page.
# (Citations in the main response.)
OPENAI_TTS_ENDPOINT = "https://api.openai.com/v1/audio/speech"

async def tts_stream(text: str, voice: str | None = None,
audio_format: str | None = None) -> AsyncIterator[bytes]:
    payload = {
        "model": settings.openai_tts_model,
```

```

    "input": text,
    # "stream": True    # If/when OpenAI supports streaming
responses for TTS in this endpoint.
    "voice": voice or "alloy",
    "format": (audio_format or settings.audio_format),
}

headers = {
    "Authorization": f"Bearer {settings.openai_api_key}",
    "Content-Type": "application/json",
}

async with httpx.AsyncClient(timeout=None) as client:
    # We use stream=True to forward the response as it arrives.
    async with client.stream("POST", OPENAI_TTS_ENDPOINT,
headers=headers, json=payload) as resp:
        resp.raise_for_status()
        async for chunk in resp.aiter_bytes():
            if chunk:
                yield chunk

```

services/openai_summarize.py

```

import httpx
from settings import settings

# For summarization we prefer the newer Responses API semantics as
they roll out.
# Until fully standardized, this uses the /chat/completions-equivalent
JSON.
# If you migrate to the Responses API, change ENDPOINT & payload shape
accordingly.

OPENAI_TEXT_ENDPOINT = "https://api.openai.com/v1/chat/completions"

SUMMARY_SYS = (

```

```

    "You are a concise academic summarizer. "
    "Write a faithful, non-speculative abstract of the input text. "
    "Return 3-7 sentences, preserve key terms, and avoid quotations."
)

async def summarize(text: str, max_tokens: int = 300) -> str:
    headers = {
        "Authorization": f"Bearer {settings.openai_api_key}",
        "Content-Type": "application/json",
    }
    payload = {
        "model": settings.openai_summary_model,
        "messages": [
            {"role": "system", "content": SUMMARY_SYS},
            {"role": "user", "content": text},
        ],
        "temperature": 0.2,
        "max_tokens": max_tokens,
    }

    async with httpx.AsyncClient(timeout=120) as client:
        r = await client.post(OPENAI_TEXT_ENDPOINT, headers=headers,
        json=payload)
        r.raise_for_status()
        data = r.json()
        return data["choices"][0]["message"]["content"].strip()

```

*(If you adopt the **Responses API** later, switch to its request/streaming format per docs. [OpenAI Platform](#))*

routers/tts.py

```

from fastapi import APIRouter, Depends, HTTPException, status
from fastapi.responses import StreamingResponse
from pydantic import BaseModel, Field
from typing import Optional

```

```

from auth import verify_bearer
from services.openai_tts import tts_stream
from settings import settings

router = APIRouter(prefix="/api/v1", tags=["tts"])

class TTSIn(BaseModel):
    text: str = Field(min_length=1, max_length=200_000)
    voice: Optional[str] = None
    format: Optional[str] = None

@router.post("/tts")
async def tts(req: TTSIn, _=Depends(verify_bearer)):
    if not req.text.strip():
        raise HTTPException(status_code=status.HTTP_400_BAD_REQUEST,
detail="Text is empty")
    fmt = (req.format or settings.audio_format).lower()
    media_type = {
        "mp3": "audio/mpeg",
        "wav": "audio/wav",
        "ogg": "audio/ogg",
    }.get(fmt, "audio/mpeg")

    gen = tts_stream(text=req.text, voice=req.voice, audio_format=fmt)
    return StreamingResponse(gen, media_type=media_type)

```

routers/summarize.py

```

from fastapi import APIRouter, Depends, HTTPException, status
from pydantic import BaseModel, Field
from auth import verify_bearer
from services.openai_summarize import summarize

router = APIRouter(prefix="/api/v1", tags=["summarize"])

class SummarizeIn(BaseModel):

```

```
text: str = Field(min_length=1, max_length=200_000)
max_tokens: int = Field(default=300, ge=64, le=2000)

class SummarizeOut(BaseModel):
    summary: str

@router.post("/summarize", response_model=SummarizeOut)
async def summarize_endpoint(req: SummarizeIn,
                            _=Depends	verify_bearer):
    if not req.text.strip():
        raise HTTPException(status_code=status.HTTP_400_BAD_REQUEST,
                            detail="Text is empty")
    out = await summarize(req.text, max_tokens=req.max_tokens)
    return SummarizeOut(summary=out)
```

app.py

```
from fastapi import FastAPI
from routers import tts, summarize

app = FastAPI(title="ReadMe Cloud API", version="0.1.0")
app.include_router(tts.router)
app.include_router(summarize.router)

@app.get("/api/v1/health")
async def health():
    return {"ok": True}
```

Minimal tests (pytest-style)

tests/test_tts.py

```
import os
```

```
from httpx import AsyncClient
from app import app

API_BEARER = os.getenv("API_BEARER", "test-token")

async def test_tts_health():
    async with AsyncClient(app=app, base_url="http://test") as ac:
        r = await ac.get("/api/v1/health")
        assert r.status_code == 200

async def test_tts_requires_auth():
    async with AsyncClient(app=app, base_url="http://test") as ac:
        r = await ac.post("/api/v1/tts", json={"text": "Hello"})
        assert r.status_code == 401
```

tests/test_summarize.py

```
import os
from httpx import AsyncClient
from app import app

API_BEARER = os.getenv("API_BEARER", "test-token")

async def test_summarize_requires_auth():
    async with AsyncClient(app=app, base_url="http://test") as ac:
        r = await ac.post("/api/v1/summarize", json={"text": "x"*100})
        assert r.status_code == 401
```

Example calls from your local app

TTS (streamed MP3):

```
curl -X POST https://<your-heroku-app>.herokuapp.com/api/v1/tts \
-H "Authorization: Bearer $API_BEARER" \
-H "Content-Type: application/json" \
```

```
--data '{"text":"Chapter one begins...", "voice":"alloy",  
"format":"mp3"}' \  
--output chapter1.mp3
```

Summarize:

```
curl -X POST https://<your-heroku-app>.herokuapp.com/api/v1/summarize  
\  
-H "Authorization: Bearer $API_BEARER" \  
-H "Content-Type: application/json" \  
--data '{"text":"<long passage here>", "max_tokens": 300}'
```

Security and durability notes

- Endpoints require your **personal bearer**; rotate **API_BEARER** regularly.
 - All OpenAI calls are HTTPS and scoped by **OPENAI_API_KEY**.
 - No files persist on Heroku; you stream bytes directly back to the desktop client.
 - If you later add job queues or pre-generation, introduce a signed, time-limited URL layer.
-

Where this aligns with current OpenAI docs

- **Text-to-Speech**: Audio API “audio/speech” with models incl. **gpt-4o-mini-tts**; guide and model pages outline inputs and voice options. [OpenAI Platform+2](#)
[OpenAI Platform+2](#)
- **Model/version hygiene**: Monitor **Changelog** and **Deprecations**; if you switch to the **Responses API** for text, adjust the summarization service accordingly. [OpenAI Platform+2](#)
[OpenAI Platform+2](#)

